

```
import numpy as np
import tensorflow as tf
import matplotlib.pyplot as plt
from tabulate import tabulate
layers = tf.keras.layers
```

The code block below defines a few helper functions to visualize the results. You do not need to touch them.

```
def plot_examples(X, Y, n=10):
    """ Plot the first n examples for each of the 10 classes in the CIFAR dataset X, Y """
    fig, axes = plt.subplots(n, 10, figsize=(10, n))
    for l in range(10):
        axes[0, l].set_title(cifar10_labels[l], fontsize="smaller")
        m = np.squeeze(Y) == l # boolean mask: True for all images of label l
        for i in range(n):
            image = X[m][i].astype("uint8") # imshow expects uint8
            ax = axes[i, l]
            ax.imshow(image, origin="upper")
            ax.set(xticks=[], yticks=[])
    return fig, ax

def plot_prediction(X, Y, Y_predict):
    """
    Plot image X along with predicted probabilities Y_predict.
    X: CIFAR image, shape = (32, 32, 3)
    Y: CIFAR label, one-hot encoded, shape = (10)
    Y_predict: predicted probabilities, shape = (10)
    """
    fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(8, 4))

    # plot image
    ax1.imshow(X.astype("uint8"), origin="upper")
    ax1.set(xticks=[], yticks=[])

    # plot probabilities
    ax2.barh(np.arange(10), Y_predict, align="center")
    ax2.set(xlim=(0, 1), xlabel="Score", yticks=[])
    for i in range(10):
        c = "red" if (i == np.argmax(Y)) else "black"
        ax2.text(0.05, i, cifar10_labels[i].capitalize(), ha="left", va="center", color=c)

def plot_confusion(Y_true, Y_predict):
    """
    Plot confusion matrix
    Y_true: array of true classifications (0-9), shape = (N)
    Y_predict: array of predicted classifications (0-9), shape = (N)
    """
```

```

C = np.histogram2d(Y_true, Y_predict, bins=np.linspace(-0.5, 9.5, 11))[0]
Cn = C / np.sum(C, axis=1)

fig = plt.figure()
plt.imshow(Cn, interpolation="nearest", vmin=0, vmax=1, cmap=plt.cm.YlGnBu)
plt.colorbar()
plt.xlabel("prediction")
plt.ylabel("truth")
plt.xticks(range(10), cifar10_labels, rotation="vertical")
plt.yticks(range(10), cifar10_labels)
for x in range(10):
    for y in range(10):
        plt.annotate("%i" % C[x, y], xy=(y, x), ha="center", va="center")

```

First we load and preprocess CIFAR-10 data. The images are 32x32 pixels and have three color channels (red, green blue).

```

# X: images, Y: labels
(x_train, y_train), (x_test, y_test) = tf.keras.datasets.cifar10.load_data()

print("images, shape = ", x_train.shape)
print("labels, shape = ", y_train.shape)

```

```

cifar10_labels = np.array([
    'airplane',
    'automobile',
    'bird',
    'cat',
    'deer',
    'dog',
    'frog',
    'horse',
    'ship',
    'truck'])

```

⏏ Downloading data from <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>
 170498071/170498071 [=====] - 9s 0us/step
 images, shape = (50000, 32, 32, 3)
 labels, shape = (50000, 1)

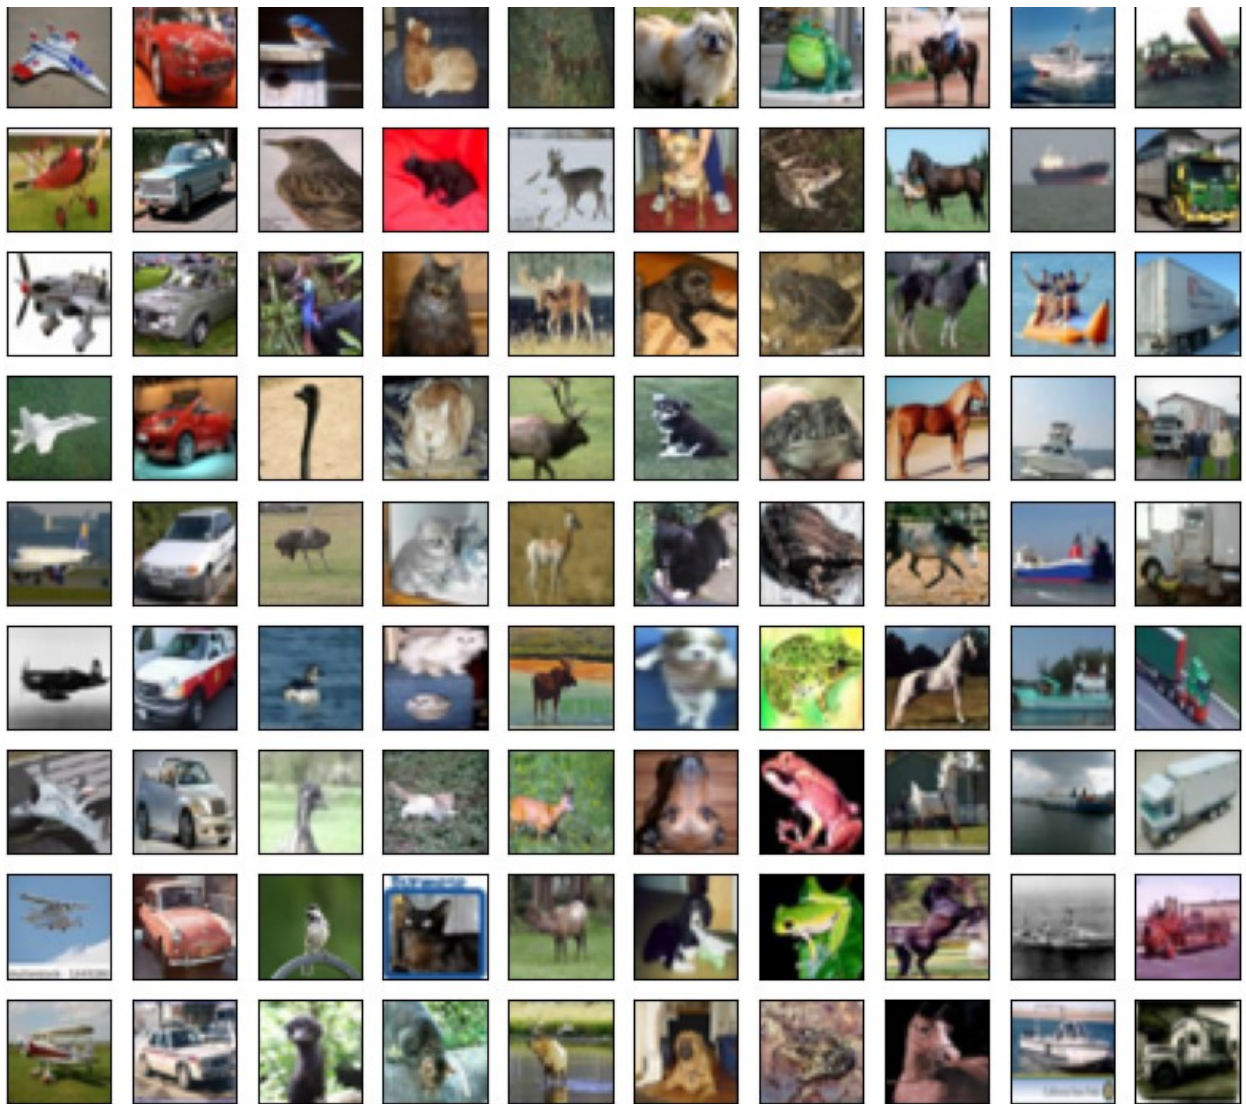
```

# Hint: To plot example images, you can use the plot examples function
plot_examples(x_train, y_train)

```

⏏ (<Figure size 1000x1000 with 100 Axes>, <Axes: >)

airplane	automobile	bird	cat	deer	dog	frog	horse	ship	truck
									



```
print(y_test.shape)
```

```
(10000, 1)
```

```
# convert labels ("0"-"9") to one-hot encodings, "0" = (1, 0, ... 0) and so on
y_train_onehot = tf.keras.utils.to_categorical(y_train, 10)
y_test_onehot = tf.keras.utils.to_categorical(y_test, 10)[:8000]
y_valid_onehot = tf.keras.utils.to_categorical(y_test, 10)[8000:]
```

```
# Hint: normalize the data
```

```
x_train_norm = (x_train-np.mean(x_train)) / np.std(x_train)
x_test_norm = (x_test-np.mean(x_test)) / np.std(x_test)
```

```
# Hint: use 20% of the training data for validation
```

```

x_valid_norm=x_test_norm[8000:]
x_test_norm=x_test_norm[:8000]

```

▼ Define dense network

```

def dense_block(x,n=4, channels=8, convshape=(3,3)):
    xlist=[x]
    for i in range(n):
        xlist.append(layers.Conv2D(channels, convshape, padding="same", activation="relu"
        x=layers.concatenate(xlist[:], axis=-1)
    return x

#Implement the horse example from slide 17 in the notes
x0=tf.keras.Input(shape=(32, 32, 3), name="input")
x=layers.Conv2D(8, kernel_size=(3, 3), padding="valid", activation="relu")(x0)
x=dense_block(x, 4, 16, (4,4))
x=layers.Conv2D(16, kernel_size=(4, 4), padding="valid", activation="relu")(x)
x=layers.MaxPooling2D((2,2), strides=(2,2))(x)

x=dense_block(x, 4, 32, (3,3))
x=layers.Conv2D(32, kernel_size=(3, 3), padding="same", activation="relu")(x)
x=layers.MaxPooling2D((3,3), strides=(3,3))(x)

x=dense_block(x, 4, 64, (2,2))
x=layers.MaxPooling2D((2,2), strides=(2,2))(x)
x=layers.Flatten()(x)
x=layers.Dense(10, activation="softmax")(x)

model = tf.keras.models.Model(
    inputs=[x0],
    outputs=[x],
    name="DenseNet",
)

print(model.summary())

```

Model: "DenseNet"

Layer (type)	Output Shape	Param #	Connected to
input (InputLayer)	[(None, 32, 32, 3)]	0	[]
conv2d (Conv2D)	(None, 30, 30, 8)	224	['input[0][0]']
conv2d_1 (Conv2D)	(None, 30, 30, 16)	2064	['conv2d[0][0]']
concatenate (Concatenate)	(None, 30, 30, 24)	0	['conv2d[0][0]', 'conv2d_1[0][0]']
conv2d_2 (Conv2D)	(None, 30, 30, 16)	6160	['concatenate[0][0]']
concatenate_1 (Concatenate)	(None, 30, 30, 40)	0	['conv2d[0][0]', 'conv2d_2[0][0]']

```

conv2d_1 (Conv2D)          (None, 30, 30, 16)  10256  ['conv2d_1[0][0]']
conv2d_2 (Conv2D)          (None, 30, 30, 16)  10256  ['conv2d_2[0][0]']

conv2d_3 (Conv2D)          (None, 30, 30, 16)  10256  ['concatenate_1[0]']

concatenate_2 (Concatenate) (None, 30, 30, 56)   0        ['conv2d[0][0]',
        'conv2d_1[0][0]',
        'conv2d_2[0][0]',
        'conv2d_3[0][0]']

conv2d_4 (Conv2D)          (None, 30, 30, 16)  14352  ['concatenate_2[0]']

concatenate_3 (Concatenate) (None, 30, 30, 72)   0        ['conv2d[0][0]',
        'conv2d_1[0][0]',
        'conv2d_2[0][0]',
        'conv2d_3[0][0]',
        'conv2d_4[0][0]']

conv2d_5 (Conv2D)          (None, 27, 27, 16)  18448  ['concatenate_3[0]']

max_pooling2d (MaxPooling2D) (None, 13, 13, 16)   0        ['conv2d_5[0][0]']

conv2d_6 (Conv2D)          (None, 13, 13, 32)  4640    ['max_pooling2d[0]']

concatenate_4 (Concatenate) (None, 13, 13, 48)   0        ['max_pooling2d[0]',
        'conv2d_6[0][0]']

conv2d_7 (Conv2D)          (None, 13, 13, 32)  13856  ['concatenate_4[0]']

concatenate_5 (Concatenate) (None, 13, 13, 80)   0        ['max_pooling2d[0]',
        'conv2d_6[0][0]',
        'conv2d_7[0][0]']

conv2d_8 (Conv2D)          (None, 13, 13, 32)  23072  ['concatenate_5[0]']

concatenate_6 (Concatenate) (None, 13, 13, 112)  0        ['max_pooling2d[0]',
        'conv2d_6[0][0]',
        'conv2d_7[0][0]',
        'conv2d_8[0][0]']

conv2d_9 (Conv2D)          (None, 13, 13, 32)  32288  ['concatenate_6[0]']

```

```

# -----
# Training
# -----
model.compile(
    loss='categorical_crossentropy',
    optimizer="adam",
    metrics=["accuracy"]
)

model.fit(
    x_train_norm, y_train_onehot,
    batch_size=1000,
    epochs=20, # train at least for 20 epochs
    verbose=1,
    validation_data=(x_valid_norm, y_valid_onehot).

```

```

callbacks=[tf.keras.callbacks.CSVLogger("history_{}.csv".format(model.name))],
)

Epoch 1/20
50/50 [=====] - 49s 719ms/step - loss: 1.9764 - accuracy:
Epoch 2/20
50/50 [=====] - 13s 256ms/step - loss: 1.5586 - accuracy:
Epoch 3/20
50/50 [=====] - 13s 253ms/step - loss: 1.3427 - accuracy:
Epoch 4/20
50/50 [=====] - 13s 252ms/step - loss: 1.2093 - accuracy:
Epoch 5/20
50/50 [=====] - 12s 246ms/step - loss: 1.1043 - accuracy:
Epoch 6/20
50/50 [=====] - 12s 246ms/step - loss: 1.0235 - accuracy:
Epoch 7/20
50/50 [=====] - 12s 248ms/step - loss: 0.9392 - accuracy:
Epoch 8/20
50/50 [=====] - 12s 248ms/step - loss: 0.8757 - accuracy:
Epoch 9/20
50/50 [=====] - 12s 248ms/step - loss: 0.8196 - accuracy:
Epoch 10/20
50/50 [=====] - 12s 248ms/step - loss: 0.7555 - accuracy:
Epoch 11/20
50/50 [=====] - 12s 249ms/step - loss: 0.7115 - accuracy:
Epoch 12/20
50/50 [=====] - 12s 248ms/step - loss: 0.6614 - accuracy:
Epoch 13/20
50/50 [=====] - 12s 249ms/step - loss: 0.6224 - accuracy:
Epoch 14/20
50/50 [=====] - 15s 310ms/step - loss: 0.5737 - accuracy:
Epoch 15/20
50/50 [=====] - 31s 614ms/step - loss: 0.5397 - accuracy:
Epoch 16/20
50/50 [=====] - 13s 251ms/step - loss: 0.4854 - accuracy:
Epoch 17/20
50/50 [=====] - 12s 248ms/step - loss: 0.4716 - accuracy:
Epoch 18/20
50/50 [=====] - 12s 248ms/step - loss: 0.4265 - accuracy:
Epoch 19/20
50/50 [=====] - 12s 249ms/step - loss: 0.3676 - accuracy:
Epoch 20/20
50/50 [=====] - 12s 249ms/step - loss: 0.3656 - accuracy:
<keras.callbacks.History at 0x222cc0a1250>

```

```

# -----
# Plots
# -----
# training curves
history = np.genfromtxt("history_{}.csv".format(model.name), delimiter=",", names=True

plt.figure()
plt.plot(history["epoch"], history["accuracy"], label="training accuracy")
plt.plot(history["epoch"], history["val_accuracy"], label="validation accuracy")
plt.legend()
plt.xlabel("Epoch")
plt.ylabel("Accuracy")

```

```

plt.ylabel('Accuracy')

# Hint: this is how you can plot the confusion matrix.
# calculate predictions for test set
y_predict = model.predict(x_test_norm, batch_size=128)

# convert back to class labels (0-9)
y_predict_cl = np.argmax(y_predict, axis=1)
y_test_cl = np.argmax(y_test_onehot, axis=1)

# plot confusion matrix
plot_confusion(y_test_cl, y_predict_cl)

#
test_acc=np.mean(np.equal(y_predict_cl, y_test_cl))
print(test_acc)

# Task: plot a few examples of correctly and incorrectly classified images.
# Hint: First find the indices of correctly and incorrectly classified images:
i0 = np.argwhere(y_predict_cl != y_test_cl) # misclassified images
i1 = np.argwhere(y_predict_cl == y_test_cl) # correctly classified images


# original (unnormalized) test images
x_test = x_test[:8000]

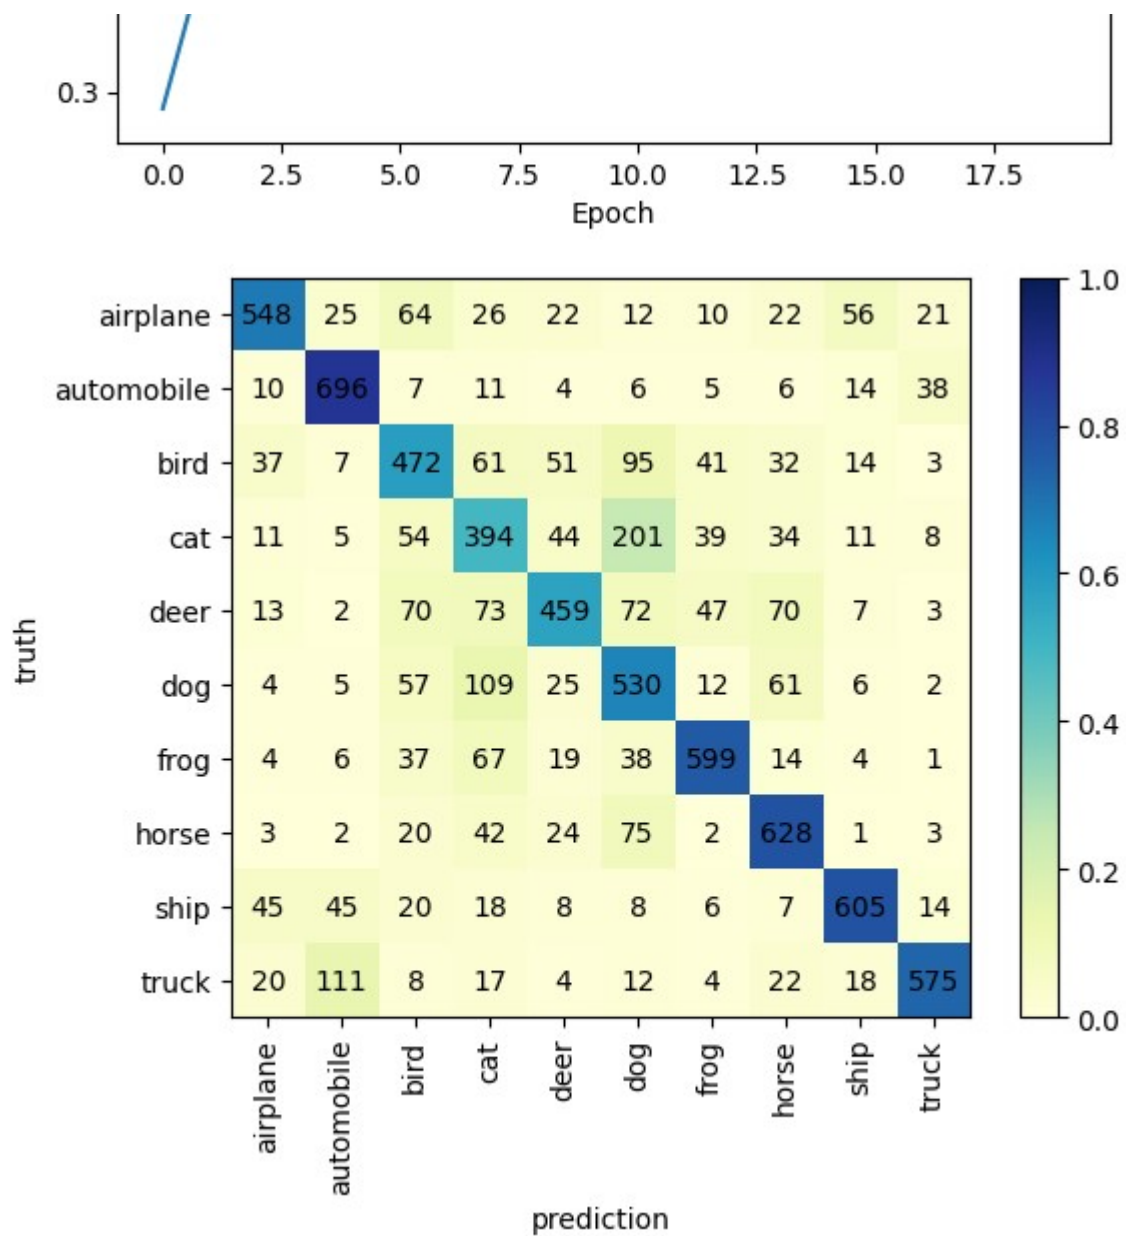
# Hint: Now you can use the `plot_prediction` function to plot the images:
# plot first 3 false classifications
for i in i0[0:3,0]:
    plot_prediction(x_test[i], y_test_onehot[i], y_predict[i])

#Plot correct classifications:
for i in i1[0:3,0]:
    plot_prediction(x_test[i], y_test_onehot[i], y_predict[i])

```


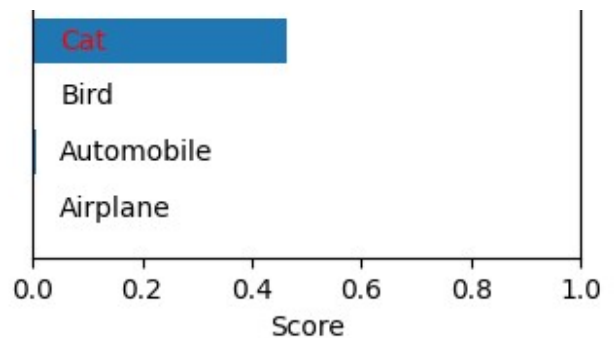

63/63 [=====] - 4s 47ms/step
0.68825






Note

train accuracy: 87%, validation accuracy: 69%, test accuracy: 72%. Much better test accuracy now=better model with DenseNet. Although it is a bit overfitted (high generalization error), we could use some regularization or similar to improve performance.



Truck



Ship

Horse

Frog

Dog

Deer

Cat

Bird